



INSTITUTE FOR DEFENSE ANALYSES

Power Approximations for Reliability Test Designs

Yevgeniya Pinelis, *Project Leader*

Rebecca Medlin
Laura Freeman
Thomas Johnson

September 2018

Approved for public release.
Distribution is unlimited.

IDA Non-Standard Document
NS D-9183

Log: H 2018-000406

INSTITUTE FOR DEFENSE ANALYSES
4850 Mark Center Drive
Alexandria, Virginia 22311-1882



The Institute for Defense Analyses is a non-profit corporation that operates three federally funded research and development centers to provide objective analyses of national security issues, particularly those requiring scientific and technical expertise, and conduct related research on other national challenges.

About This Publication

This work was conducted by the Institute for Defense Analyses (IDA) under contract HQ0034-14-D-0001, Task AE C9082, "Statistics and Data Science WG," for the Office of the Director, Operational Test and Evaluation. The views, opinions, and findings should not be construed as representing the official position of either the Department of Defense or the sponsoring organization.

Acknowledgments

The IDA Technical Review Committee was chaired by Mr. Robert R. Soule and consisted of Dr. John Haman from the Operational Evaluation Division.

For more information:

Yevgeniya Pinelis, Project Leader
ypinelis@ida.org • 703-845-6899

Robert R. Soule, Director, Operational Evaluation Division
rsoule@ida.org • (703) 845-2482

Copyright Notice

© 2018 Institute for Defense Analyses
4850 Mark Center Drive, Alexandria, Virginia 22311-1882 • (703) 845-2000.

This material may be reproduced by or for the U.S. Government pursuant to the copyright license under the clause at DFARS 252.227-7013 (a)(16) [Jun 2013].

INSTITUTE FOR DEFENSE ANALYSES

IDA Non-Standard Document NS D-9183

Power Approximations for Reliability Test Designs

Yevgeniya Pinelis, *Project Leader*

Rebecca Medlin
Laura Freeman
Thomas Johnson

Executive Summary

Reliability tests determine which factors drive system reliability. Often, the reliability or failure time data collected in these tests tend to follow distinctly non-normal distributions and include censored observations. The experimental design should accommodate the skewed nature of the response and allow for censored observations, which occur when systems under test do not fail within the allotted test time. To account for these design and analysis considerations, Monte Carlo simulations are frequently used to evaluate experimental design properties. Simulation provides accurate power calculations as a function of sample size, and allows researchers to determine adequate sample sizes at each level of the treatment. However, simulation may be inefficient to compare multiple experiments of various sizes.

In this document, we present a closed form approach to calculate power, based on the non-central chi-squared approximation to the distribution of the likelihood ratio statistic. The solution can be used to compare multiple designs and accommodate trade-space analyses between power, effect size, model formulation, sample size, censoring rates, and design

type. To demonstrate the efficiency of our approach, we provide a comparison to estimates that are generated using Monte Carlo simulation.

At present, our power approximation method to evaluate the power associated with a reliability test design is available to IDA analysts at <https://test-science.shinyapps.io/survpow/>.



Power Approximations for Reliability Test Designs

Tom Johnson
Rebecca Medlin
Laura Freeman

October 2018

Synopsis

Reliability experiments determine which factors drive system/product reliability.

Reliability data tend to follow distinctly non-normal distributions and include censored observations.

Our experimental designs should accommodate the skewed nature of the response and allow for censored observations.

Monte Carlo simulations are frequently used to evaluate the design properties (e.g., power) for reliability experiments.

Simulation can be inefficient to compare multiple experiments of various sizes.

We have developed a closed form approximation to calculate the power of a reliability experiment.

Illustrative Example

Planning a reliability experiment for $N = 240$ “products.”

A $2 \times 2 \times 3$ full factorial experiment (**design**)

Expose until failure or up to **10 days** (**right censoring time**)

Identify seven days as the crucial juncture in the product’s lifetime; we anticipate 80% of the products will fail by this time (**nominal failure rate**).

Detect 10% change in probability of failure due to an exposure factor (**effect size**).

Lower probability of failure $p_1 = .75$ and upper probability of failure $p_2 = .85$

Failure times follow a **lognormal** distribution with fixed scale parameter, $\sigma = 2$

Effect sizes, in terms of the location parameters μ_1 and μ_2 are found by:

$$p_1 = F(t_p, \mu_{p_1}) \text{ and } p_2 = F(t_p, \mu_{p_2}):$$

$$\mu_1 = .6 \text{ and } \mu_2 = -.13$$

How do we evaluate our design?

How do we calculate the **power of test** for our nine model coefficients?

Discussion of power is common in classical experimental design evaluation

Discussions of power **are not** prevalent in reliability research.

Meeker (1977; 1992; 1992; 1994; 1995; 1998; 2006).

- Precision around a Quantile Estimate or Hazard Functions
- Good for quality control applications...

So why not just use Monte Carlo?

It is a flexible and accurate approach.

It could quickly become computationally inefficient.

A closed form approximation is computationally efficient.

The Failure-Time Regression Model

$$y_i = \log(T_i) = m_i^T \beta + \sigma \epsilon_i$$

$i = 1, 2, \dots, k$ design points

Model the location $\mu_i = m_i^T \beta$; fix the scale σ

Lognormal Model: $T_i \sim \text{LogN}(\mu_i, \sigma)$ and $\epsilon_i \sim N(0, 1)$

Weibull Model: $T_i \sim \text{Weibull}(\mu_i, \sigma)$ and $\epsilon_i \sim \text{SEV}(0, 1)$

Censoring: Fixed, Type I Right Censoring Scheme

$$\delta_{ij} = 1 \text{ if } T_{ij} < t_c ; \delta_{ij} = 0 \text{ otherwise}$$

Maximum likelihood estimation

$$l_{n\sigma} = \sum_{i=1}^k \sum_{j=1}^{n_i} \log [f_{T_{ij}, \mu_i}] \delta_{ij} + \log[1 - F_{t_c, \mu_i}](1 - \delta_{ij})$$

Testing for Significance: Likelihood Ratio Test

Coefficients under test

Nuisance coefficients

$$\beta = [\psi, \lambda] \text{ and } M = [X, Z]$$

$$\mu_i = x_i^T \psi + z_i^T \lambda$$


$$H_0: \psi = \psi_0$$

$$H_A: \psi \neq \psi_0$$

$$\text{Likelihood Ratio Statistic: } 2[l_{n\sigma}(\hat{\psi}, \hat{\lambda}) - l_{n\sigma}(\psi_0, \hat{\lambda}_0)]$$

The power of a test is:

$$Pr(t > X_{\alpha}^2) = 1 - \tilde{X}^2(X_{\alpha}^2, p, \gamma)$$



Central χ^2 and non-central \tilde{X}^2

t is the non-central chi-square random variable.

X_{α}^2 is the upper α percentage point of the central chi-square distribution

\tilde{X}^2 is the non-central chi-square distribution with p degrees of freedom (number of coefficients under test) and non-centrality parameter γ

Self, Mauritsen, and O'Hara (1992)

Describe a power approximation approach for the generalized linear model (glm) framework.

Based on a non-central chi-square approximation to the distribution of the likelihood ratio statistic.

$$LRS \sim \tilde{X}_{p,\gamma}^2$$

Technique accommodates any model within the exponential family of distributions that can be arranged into the glm canonical form:

$$f(y_i) = \exp \left(\frac{y_i \theta_i - b(\theta_i)}{a_i(\phi)} + c(y_i, \phi) \right)$$

Examples: logistic, Poisson, and gamma regression models.

Failure-time regression models share many qualities of the generalized linear model, but they cannot be arranged into the canonical form.

If we can solve for the non-centrality parameter γ , then we can estimate power.

We know the expected value of a noncentral chi-square random variable is $p + \gamma$.

If we equate the expected value of the LRS to the expected value of the noncentral chi-square random variable, we can solve for γ .

Great Expectations [and Expansions]

Coefficients under test

Nuisance coefficients

$$\begin{aligned} E_{\psi, \lambda} \{ 2[l_{n\sigma}(\hat{\psi}, \hat{\lambda}) - l_{n\sigma}(\psi_0, \hat{\lambda}_0)] \} = \\ E_{\psi, \lambda} \{ 2[l_{n\sigma}(\hat{\psi}, \hat{\lambda}) - l_{n\sigma}(\psi, \lambda)] \} \\ - E_{\psi, \lambda} \{ 2[l_{n\sigma}(\psi_0, \hat{\lambda}_0) - l_{n\sigma}(\psi_0, \lambda_0^*)] \} \\ + E_{\psi, \lambda} \{ 2[l_{n\sigma}(\psi, \lambda) - l_{n\sigma}(\psi_0, \lambda_0^*)] \} \end{aligned}$$

$$p + \gamma = A - B + C$$

$E_{\psi, \lambda} \{ \cdot \}$ taken with respect to the true parameters ψ and λ

λ_0^* is the limiting value of the null coefficients

$$p + \gamma = A - B + C$$

The first term (A): $A = p + q$

- Where p is the number of coefficients under test and q is the number of nuisance coefficients (i.e., coefficients not under test)

The second term (B): take limits and use Taylor series expansion...

- Closed form solution for GLM within the canonical form (Self et al., 1992).
- **No closed form solution for failure-time regression models with censoring.**

A numerical solution is possible, but would undermine our objective!

The third term (C): a closed form solution exists!

- Closed form solution for GLM (shown in Self et al., 1992)
- **Closed form solution for Failure Time regression models with fixed, Type 1, right-censored data (our work)**

All you need is C

$$\gamma = A - B + C - p$$
$$\approx$$
$$\gamma = C$$

“the dominance of the [C] term in the calculation of the non-centrality parameter,” - Self et al. (1992)

“In our experience, the term [A – B – p] is usually very close to zero.” - Self et al. (1992), O’Brien and Shieh (1998), Shieh (2000), Brown et al. (1999)

Solving for C

$$C = E_{\psi, \lambda} \{ \log [f_{T_{ij}, \mu_i}] \delta_{ij} + \log [1 - F_{t_c, \mu_i}] (1 - \delta_{ij}) - \log [f_{T_{ij}, \mu_i^*}] \delta_{ij} + \log [1 - F_{t_c, \mu_i^*}] (1 - \delta_{ij}) \}$$

Lognormal Solution

$$C = \sum_{i=1}^k n_i \left\{ \frac{2f_{t_c, -\mu_i}(\mu_i^* - \mu_i)}{t_c^{-1 - \frac{2\mu_i}{\sigma^2}}} + \frac{F_{t_c, \mu_i}(\mu_i - \mu_i^*)^2}{\sigma^2} + (2 - 2F_{t_c, \mu_i}) \log \left(\frac{2 - 2F_{t_c, \mu_i}}{2 - 2F_{t_c, \mu_i^*}} \right) \right\}$$

Weibull Solution

$$C = \sum_{i=1}^k n_i \left\{ -\frac{2}{\sigma} e^{-\frac{\mu_i^*}{\sigma}} F_{t_c, \mu_i} \left(e^{-\frac{\mu_i}{\sigma}} \sigma + e^{\frac{\mu_i^*}{\sigma}} (\mu_i - \mu_i^* + \sigma) \right) \right\}$$

Equation Notes:

$f_{t, \mu} = f(t, \mu, \sigma)$; $F_{t, \mu} = F(t, \mu, \sigma)$

λ_0^* is the limiting value of the null coefficients and is found by fitting a lognormal regression model to the alternative data. The alternative data are the failure times that represent the perfect fit to the alternative coefficients: $T_{ij}^* = e^{\mu_i}$. Use standard failure-time model fitting software to fit the reduced model to T_{ij}^* .

The fitted coefficients are equal to λ_0^* and $\mu_i^* = x_i^T \lambda_0^* + z_i^T \psi_0$.

Returning to the Example

N: 204
 Design: 2 x 2 x 3
 Censor time: $t_c: 10$
 Nominal failure rate at 7 days: $\bar{p} = 0.8$
 Effect size: $\Delta = .1$
 $p_1 = .75$ and $p_2 = .85$
 Lognormal distribution with $\sigma = 2$
 $p_1 = F_{t_p, \mu_{p_1}}$ and $p_2 = F_{t_p, \mu_{p_1}}$
 $\mu_1 = .6$ and $\mu_2 = -.13$



Main Effects Model:
 $\beta_{int} + \beta_{F1} + \beta_{F2} + \beta_{F3_1} + \beta_{F3_2}$
 Sum-to-zero contrast, effect size of
 alternative coefficients:
 $\beta^T = [.23, .36, .36, .36, .00]$
 Illustrate power for F_1
 $H_0: \psi = 0$
 $H_A: \psi \neq 0$
 $\lambda^T = [\beta_{int}, \beta_{F2}, \beta_{F3}], \quad \psi^T = \beta_{F1}$



Power
 $Pr(t > X_{\alpha=0.05}^2)$
 $= 1 - \tilde{X}^2\{X_{\alpha=0.05}^2, p = 1, \gamma = 7.59\}$
 $= 1 - \tilde{X}^2\{3.84, p = 1, \gamma = 7.59\}$
 $= 0.7869$



Calculate C
 $\mu_i = x_i^T \lambda + z_i^T \psi$
 $\mu_i^* = x_i^T \lambda_0^*$
 $\gamma = C = 7.59$

Does the simplifying assumption work?

Model	Design	N	\bar{p}	Δ_p	σ
Lognormal	2^2	64	0.5	.06	.5
Weibull	2^6	128	0.8	.12	1
		256			3
Conditions varied using a 40-run D-optimal experimental design					
$t_c = 10; t_p = 7$					




Power Approximation (\mathcal{C})	Monte Carlo Power*
--	--------------------

Questions of Interest:


- Approximate Power: do we get a reasonable solution?
- $|\text{Approximate Power} - \text{Monte Carlo Power}|$: how accurate is that solution?
 - Is calculating \mathcal{C} good enough?

Approximate power results behave as expected



Differences from Monte Carlo Power are small

Most influential conditions: \bar{p} , Design, Model



Is the method computationally efficient?

We explore accuracy/timeliness trade-off:

- Study based on example construct.
- Power calculated for one coefficient, β_{F_1}
- 'True' power, Monte Carlo Simulation with 5 million iterations: 0.78820

Monte Carlo: Stochastic

Iteration cases considered: 100, 500, 1,000, 5,000, 10,000, 15,000, 20,000, 25,000

*Replicate each iteration case 20 times to estimate an average iteration power and corresponding accuracy, RMSE, and an average iteration time.

Power Approximation: Deterministic

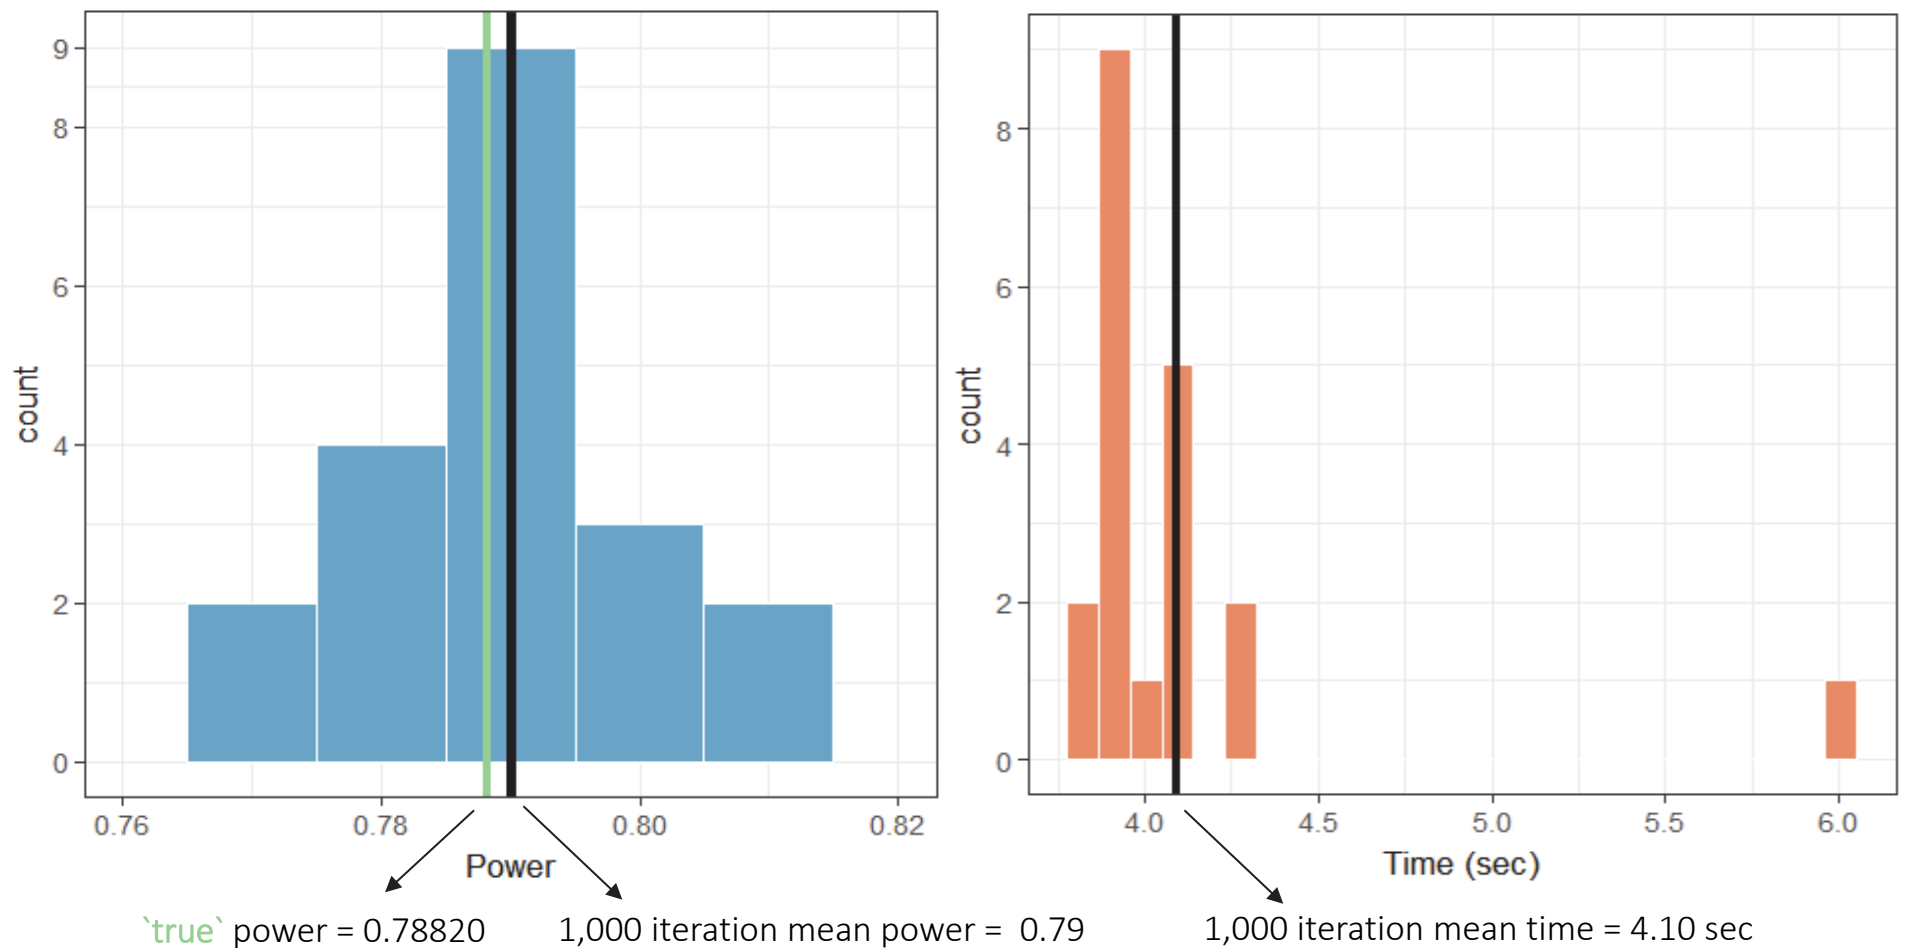
Power: 0.78688

Accuracy: 0.00132


Time: 0.25 seconds

Example Monte Carlo Output

Monte Carlo output for the 1,000 iteration case, with 20 replicates



The benefit of our approximation method relies on its computational efficiency.



Power Approximation Method

Absolute Error = .001

Time (sec) = 0.25

Computational time can quickly compound!

A main effects + two-factor interaction model has nine coefficients.

→ We need nine power computations.

Typically, confidence and power are fixed and we solve for sample size.

→ This requires iterative numerical techniques.

→ Suppose for example, our problem requires five iterations.

The 1,000 iteration case now takes:¹

Monte Carlo: 4 seconds x 9 coefficients x 5 iterations to optimize
= 3 minutes

Power Approx.: 0.25 seconds x 9 coefficients x 5 iterations to optimize
= 11 seconds

We've made it easy to implement!

<https://test-science.shinyapps.io/survpow/>

Add Factor

Levels in Factor 1

2

7

Levels in Factor 2

2

7

Levels in Factor 3

2

3

7

Sample Size

11

240

Model Order

MEs+2FIs

☐ Sim

Model

lognormal

Nominal Probability

0.8

Delta Probability

0.1

Sigma

2

Censor Time

7

Target Time

10

Alpha

0.05

Calculate Power

- **Model** specifies the expected distribution of failure times
- **Nominal Probability** is the nominal probability of failure at the Target Time.
- **Delta Probability** is the effect size. That is,

$$p_1 = p_{\text{nom}} - \frac{p_{\text{delta}}}{2}$$

$$p_2 = p_{\text{nom}} + \frac{p_{\text{delta}}}{2}$$

and the Location parameters μ_1 and μ_2 are calculated using p_1 and p_2

- **Sigma** is the scale parameter (constant)
- Failure times are not observed after the **Censor Time** (Fixed, Type 1, Right Censoring)
- **Alpha** is the Type I error rate.
- **Sample Size** adjusts the number of runs in the D-optimal experiment.
- The **simulate checkbox** allows the user to use Monte Carlo to calculate power for the likelihood ratio test.

Add Factor

Levels in Factor 1




Levels in Factor 2



Levels in Factor 3



Sample Size



Model Order

MEs+2FIs

☐ Sim

Model

lognormal

Nominal Probability

0.8

Delta Probability

0.1

Sigma

2

Censor Time

7

Target Time

10



Alpha

0.05

Calculate Power

Num	Term	dof	ncp	PowApprox
1	A	1	7.30	0.77
2	B	1	7.30	0.77
3	C	2	4.86	0.49
4	A*B	1	7.30	0.77
5	A*C	2	4.86	0.49
6	B*C	2	4.86	0.49

Num	Coeff	Value
1	(Intercept)	0.59
2	A1	-0.36
3	B1	-0.36
4	C1	-0.36
5	C2	0.00
6	A1:B1	-0.36
7	A1:C1	-0.36
8	A1:C2	0.00
9	B1:C1	-0.36
10	B1:C2	0.00



REPORT DOCUMENTATION PAGE					<i>Form Approved</i> OMB No. 0704-0188	
<p>The public reporting burden for this collection of information is estimated to average 1 hour per response, including the time for reviewing instructions, searching existing data sources, gathering and maintaining the data needed, and completing and reviewing the collection of information. Send comments regarding this burden estimate or any other aspect of this collection of information, including suggestions for reducing the burden, to Department of Defense, Washington Headquarters Services, Directorate for Information Operations and Reports (0704-0188), 1215 Jefferson Davis Highway, Suite 1204, Arlington, VA 22202-4302. Respondents should be aware that notwithstanding any other provision of law, no person shall be subject to any penalty for failing to comply with a collection of information if it does not display a currently valid OMB control number.</p> <p>PLEASE DO NOT RETURN YOUR FORM TO THE ABOVE ADDRESS.</p>						
1. REPORT DATE (DD-MM-YYYY) 01-10-2018		2. REPORT TYPE OED Draft			3. DATES COVERED (From - To) September - October 2018	
4. TITLE AND SUBTITLE Power Approximations for Reliability Test Designs				5a. CONTRACT NUMBER HQ0034-14-D-0001		
				5b. GRANT NUMBER		
				5c. PROGRAM ELEMENT NUMBER		
6. AUTHOR(S) Thomas H. Johnson (OED); Rebecca M. Medlin (OED); Laura J. Freeman (OED);				5d. PROJECT NUMBER BD-09-2299		
				5e. TASK NUMBER 2299-90		
				5f. WORK UNIT NUMBER		
7. PERFORMING ORGANIZATION NAME(S) AND ADDRESS(ES) Institute for Defense Analyses 4850 Mark Center Drive Alexandria, Virginia 22311-1882					8. PERFORMING ORGANIZATION REPORT NUMBER NS D-9183 H 2018-000406	
9. SPONSORING/MONITORING AGENCY NAME(S) AND ADDRESS(ES) Director, Operational Test and Evaluation The Pentagon 1700 Defense Washington, DC 20301					10. SPONSOR/MONITOR'S ACRONYM(S) DOT&E	
					11. SPONSOR/MONITOR'S REPORT NUMBER(S)	
12. DISTRIBUTION/AVAILABILITY STATEMENT Approved for public release. Distribution is unlimited.						
13. SUPPLEMENTARY NOTES Project Leader: Yevgeniya Pinelis						
14. ABSTRACT <p>Reliability tests determine which factors drive system reliability. Often, the reliability or failure time data collected in these tests tend to follow distinctly non-normal distributions and include censored observations. The experimental design should accommodate the skewed nature of the response and allow for censored observations, which occur when systems under test do not fail within the allotted test time. To account for these design and analysis considerations, Monte Carlo simulations are frequently used to evaluate experimental design properties. Simulation provides accurate power calculations as a function of sample size, allowing researchers to determine adequate sample sizes at each level of the treatment. However, simulation may be inefficient for comparing multiple...</p>						
15. SUBJECT TERMS Design for Reliability; Design of Experiments (DOE); Power						
16. SECURITY CLASSIFICATION OF:			17. LIMITATION OF ABSTRACT Unlimited	18. NUMBER OF PAGES 26	19a. NAME OF RESPONSIBLE PERSON Yevgeniya Pinelis	
a. REPORT Unclassified	b. ABSTRACT Unclassified	c. THIS PAGE Unclassified			19b. TELEPHONE NUMBER (Include area code) (703) 845-6899	

